

- Read
  - Write
  - Glob
  - Grep
- 

# DSPy Bootstrap Few-Shot Optimizer

## Goal

Automatically generate and select optimal few-shot demonstrations for your DSPy program using a teacher model.

## When to Use

- You have **10-50 labeled examples**
- Manual example selection is tedious or suboptimal
- You want demonstrations with reasoning traces
- Quick optimization without extensive compute

## Inputs

Input	Type	Description
program	dspy.Module	Your DSPy program to optimize
trainset	list[dspy.Example]	Training examples
metric	callable	Evaluation function
max_bootstrapped_demos	int	Max teacher-generated demos (default: 4)
max_labeled_demos	int	Max direct labeled demos (default: 16)

## Outputs

Output	Type	Description
compiled_program	dspy.Module	Optimized program with demos

## Workflow

### Phase 1: Setup

```
import dspy
from dspy.teleprompt import BootstrapFewShot

# Configure LMs
dspy.configure(lm=dspy.LM("openai/gpt-4o-mini"))
```

### Phase 2: Define Program and Metric

```
class QA(dspy.Module):
    def __init__(self):
        self.generate = dspy.ChainOfThought("question -> answer")

    def forward(self, question):
        return self.generate(question=question)

    def validate_answer(example, pred, trace=None):
        return example.answer.lower() in pred.answer.lower()
```

### Phase 3: Compile

```
optimizer = BootstrapFewShot(
    metric=validate_answer,
    max_bootstrapped_demos=4,
    max_labeled_demos=4,
    teacher_settings={'lm': dspy.LM("openai/gpt-4o")})
)

compiled_qa = optimizer.compile(QA(), trainset=trainset)
```

### Phase 4: Use and Save

```
# Use optimized program
result = compiled_qa(question="What is photosynthesis?")

# Save for production
compiled_qa.save("qa_optimized.json")
```

## Production Example

```

import dspy
from dspy.teleprompt import BootstrapFewShot
from dspy.evaluate import Evaluate
import logging

logging.basicConfig(level=logging.INFO)
logger = logging.getLogger(__name__)

class ProductionQA(dspy.Module):
    def __init__(self):
        self.cot = dspy.ChainOfThought("question -> answer")

    def forward(self, question: str):
        try:
            return self.cot(question=question)
        except Exception as e:
            logger.error(f"Generation failed: {e}")
            return dspy.Prediction(answer="Unable to answer")

    def robust_metric(example, pred, trace=None):
        if not pred.answer or pred.answer == "Unable to answer":
            return 0.0
        return float(example.answer.lower() in pred.answer.lower())

    def optimize_with_bootstrap(trainset, devset):
        """Full optimization pipeline with validation."""

        # Baseline
        baseline = ProductionQA()
        evaluator = Evaluate(devset=devset, metric=robust_metric, num_threads=4)
        baseline_score = evaluator(baseline)
        logger.info(f"Baseline: {baseline_score:.2%}")

        # Optimize
        optimizer = BootstrapFewShot(
            metric=robust_metric,
            max_bootstrapped_demos=4,
            max_labeled_demos=4
        )

        compiled = optimizer.compile(baseline, trainset=trainset)
        optimized_score = evaluator(compiled)
        logger.info(f"Optimized: {optimized_score:.2%}")

        if optimized_score > baseline_score:
            compiled.save("production_qa.json")
            return compiled

        logger.warning("Optimization didn't improve; keeping baseline")
        return baseline

```

## Best Practices

1. **Quality over quantity** - 10 excellent examples beat 100 noisy ones
2. **Use stronger teacher** - GPT-4 as teacher for GPT-3.5 student
3. **Validate with held-out set** - Always test on unseen data

4. **Start with 4 demos** - More isn't always better

## Limitations

---

- Requires labeled training data
- Teacher model costs can add up
- May not generalize to very different inputs
- Limited exploration compared to MIPROv2